

Towards integral equations at exascale in Rust

Timo Betcke

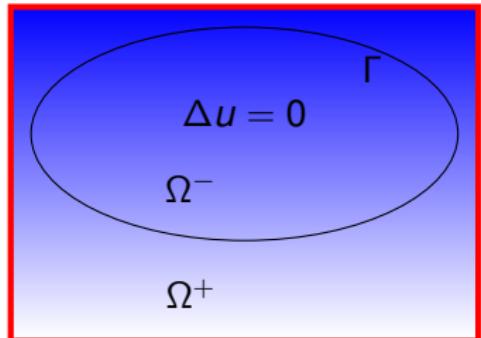
t.betcke@ucl.ac.uk

University College London

Joint with S. Kailasa, M.
Scroggs and many other
collaborators



Some Maths to set the stage



$$\Delta u(\mathbf{x}) := \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} + \frac{\partial^2 u}{\partial z^2}$$

Application areas: Heat diffusion, electrostatics, etc.

Boundary condition:
 $u(\mathbf{x}) = f(\mathbf{x}).$

Represent solution as

$$u(\mathbf{x}) = \int_{\Gamma} g(\mathbf{x}, \mathbf{y}) \phi(\mathbf{y}) ds(\mathbf{y}), \quad \mathbf{x} \in \Omega$$

Need to solve $V\phi = f$ with

$$[V\phi](\mathbf{x}) = \int_{\Gamma} g(\mathbf{x}, \mathbf{y}) \phi(\mathbf{y}) ds(\mathbf{y}), \quad \mathbf{x} \in \Gamma$$

How to discretise the integral equation?

We want to solve $V\phi = f$. Introduce a triangulation \mathcal{T} of the domain Ω into triangles τ_j . Define basis fct.

$$\phi_j(\mathbf{x}) = \begin{cases} 1 & \mathbf{x} \in \tau_j \\ 0, & \text{otherwise} \end{cases}$$

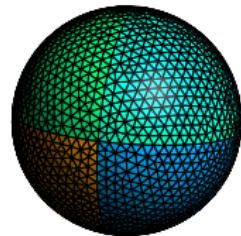
We approximate $\phi = \sum_{j=1}^N c_j \phi_j$. Multiplying with ϕ_i and integrating gives

$$\int_{\Gamma} \phi_i(\mathbf{x}) [V\phi](\mathbf{x}) d\mathbf{s}(\mathbf{x}) = \int_{\Gamma} f(\mathbf{x}) \phi_i(\mathbf{x}) d\mathbf{s}(\mathbf{x}), i = 1, \dots, N$$

Solve $\mathbf{V}\mathbf{c} = \mathbf{b}$

$$\mathbf{V}_{ij} = \int_{\tau_i} \int_{\tau_j} g(\mathbf{x}, \mathbf{y}) d\mathbf{s}(\mathbf{y}) d\mathbf{s}(\mathbf{x})$$

$$\mathbf{b}_i = \int_{\Gamma} f(\mathbf{x}) \phi_i(\mathbf{x}) d\mathbf{s}(\mathbf{x})$$



The Bempp boundary element software

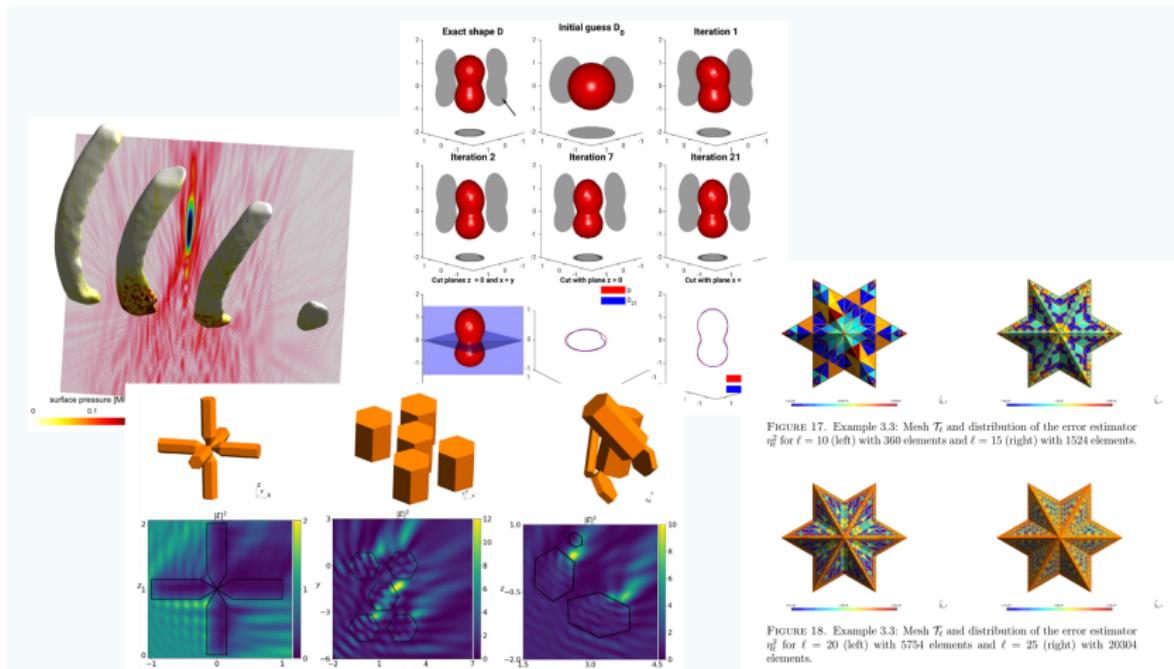


FIGURE 17. Example 3.3: Mesh T_ℓ and distribution of the error estimator η_ℓ^2 for $\ell = 10$ (left) with 360 elements and $\ell = 15$ (right) with 1524 elements.



FIGURE 18. Example 3.3: Mesh T_ℓ and distribution of the error estimator η_ℓ^2 for $\ell = 20$ (left) with 5754 elements and $\ell = 25$ (right) with 20304 elements.

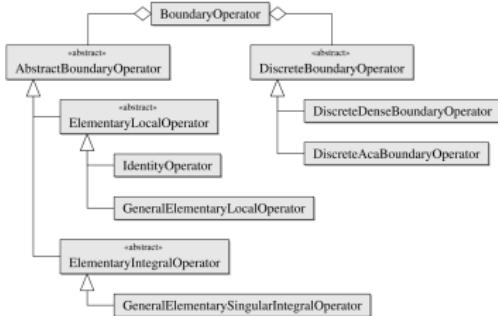


Fig. 3. Relationships between the main classes representing boundary operators.

- Complex object hierarchy
- Highly templated code
- Dependencies on complex packages (Boost, BLAS, Armadillo, Trilinos, DUNE, Intel TBB)

Python Interface

- Auto-generated with Swig
- Thin layer on C++ code
- Almost no logic in the Python layer, all handled by C++

Over time moved all high-level functions into Python, only kept low-level performance critical data structures in C++.

¹ W. Smigaj et. al., *Solving boundary integral problems with BEM++, ACM Transactions on Mathematical Software 41 (2015), pp. 1 - 40*

Numba for grid iterations

- Grid topology computations
- Python callables
- Routines on grid functions, (L^2 -norm, etc.)
- Fallback for operators

OpenCL for matrix assembly/evaluation

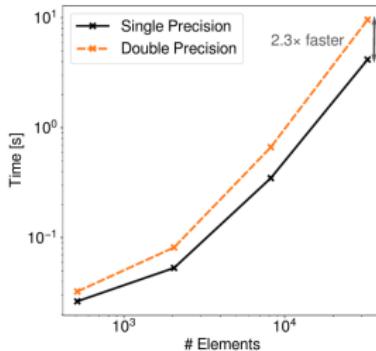
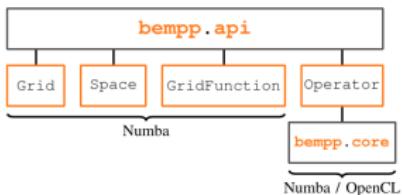
- All potential evaluations
- Mass matrices
- FMM Near-Field

Other technologies

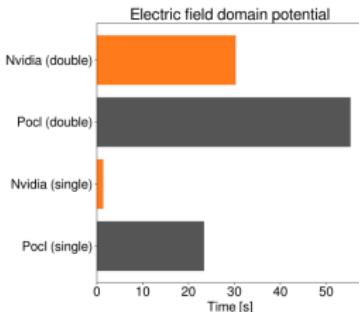
- Scipy for sparse matrix operations and iterative solvers
- Numpy dense solvers

- All logic controlled from Python (easy to hack on).
- Very fast SIMD optimized dense matrix assembly.
- But need lots of OpenCL C99 kernel code.

Characteristics of Bempp-cl²



- Top: Structure of Bempp-cl
- Top-Right: AVX acceleration in Bempp-cl
- Bottom-Right: Offloading of a domain potential operator

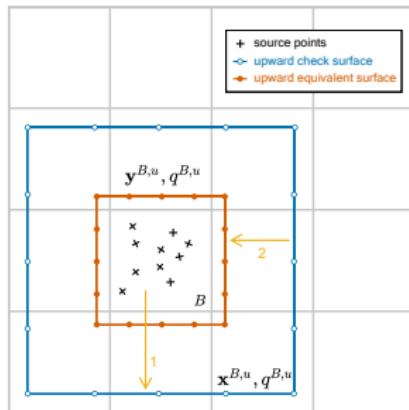
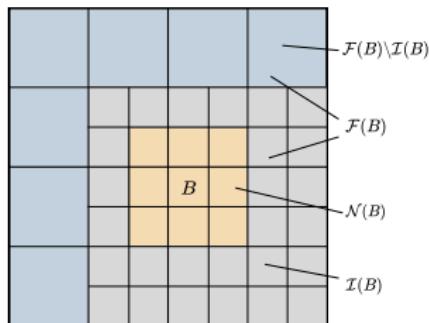
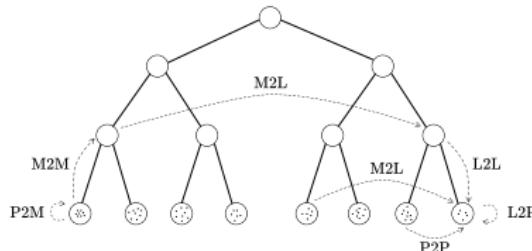


²T. Betcke, M. Scroggs, *Designing a high-performance boundary element library with OpenCL and Numba*, Computing in Science and Engineering, 23 (2021), pp. 18 - 28

Exafmm-t: High-Performance KIFMM³



- Kernel-Independent FMM library.
- Written in C++, Provides Python Interface.
- Highly efficient at higher accuracies.



³T. Wang, R. Yokota, L. Barba, *ExaFMM: a high-performance fast multipole method library with C++ and Python interfaces*, Journal of Open Source Software 6(61), 3145

Why large-scale BEM?

Solve Poisson-Boltzmann equation

$$\Delta\phi_1 = \frac{1}{\epsilon_1} \sum_k q_k \delta(\mathbf{r}, \mathbf{r}_k) \text{ in } \Omega_1$$

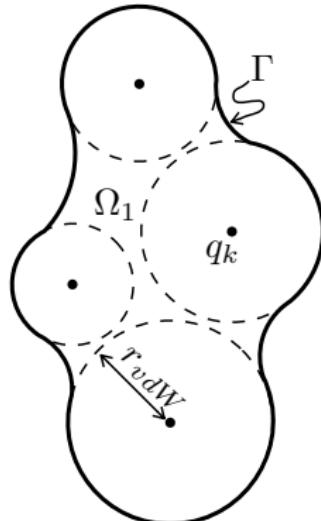
$$(\Delta - \kappa^2)\phi_2 = 0 \text{ in } \Omega_2$$

Interface conditions on Γ :

$$\phi_1 = \phi_2,$$

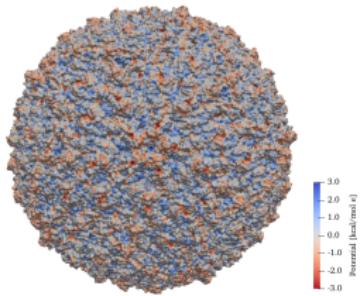
$$\epsilon_1 \frac{\partial \phi_1}{\partial \mathbf{n}} = \epsilon_2 \frac{\partial \phi_2}{\partial \mathbf{n}}$$

Let $\phi_1 = \phi_{reac} + \phi_{coul}$ (ϕ_{coul} is potential generated from the point charges only). Then want to compute the solvation energy.



$$\Delta G_{solv}^{polar} = \frac{1}{2} \sum_{k=1}^{N_q} q_k \phi_{reac}(\mathbf{r}_k)$$

Exascale FEM/BEM modeling



- Electrostatic simulation of Zika virus (10m elements) on single compute node⁴
- Mass-matrix preconditioned single trace formulation.
- FMM acceleration via ExaFMM.

Goal: Full FEM/BEM modelling of complex protein surfaces.

$$\begin{bmatrix} \epsilon_1 A & -\epsilon_1 M^T \\ (\frac{1}{2}I + K) & \frac{\epsilon_1}{\epsilon_2} V \end{bmatrix} \begin{bmatrix} \vec{\phi}_1 \\ \vec{\lambda}_2 \end{bmatrix} = \begin{bmatrix} \vec{f} \\ 0 \end{bmatrix}.$$

A: FEM matrix, K : Double-Layer operator, V : Single-Layer operator,
 M : Mass matrix

Want to scale up to billions surface dofs.

⁴T. Wang, C. D. Cooper, T. Betcke, L. A. Barba, *High-productivity, high-performance workflow for virus-scale electrostatic simulations with Bempp-Exafmm*, submitted (2021)

Deconstructing a BEM operator call

Discrete integral operator evaluation takes the form

$$Vx = P_{dual}^T(G - C)P_{dom}\bar{x} + S\bar{x}$$

- G dense.
- All other matrices highly sparse.

- $P_{dom/dual}$: Map from domain/dual space coefficients to values at quadrature points
 - G matrix of Green's fct. evaluations at regular Gauss points
 - C : Correction matrix containing contribution of G associated with adjacent triangles
 - S : Contributions of singular quadrature rule on adjacent triangles
-
- Parallel fast methods (H -matrices/FMM) for G .
 - Distributed sparse linear algebra for everything else.

Python

- Good for higher level logic.
- Efficient libs for array processing.
- Performance for complex algorithms in Python (e.g. FMM) difficult⁵.

Python + Code Gen

- Most flexible solution.
- Complex framework.
- What low-level language to use?

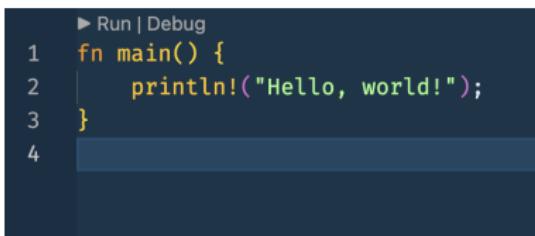
C++

- Complex Makefiles.
- Deployment difficult.
- C++20 support still patchy.

⁵S. Kailas et. al., *PyExaFMM: an exercise in designing high-performance software with Python and Numba*, to be submitted

Why Rust?

- Modern language design.
- Great tooling.
- Compile time safety guarantees.
- Simple deployment.
- Python integration trivial.



A screenshot of a code editor showing a simple Rust program. The code consists of four lines: a comment 'Run | Debug', a function definition 'fn main() {', a call to 'println!("Hello, world!");', and a closing brace '}'. The code is highlighted in yellow and purple, and the background of the editor is dark blue.

```
▶ Run | Debug
1 fn main() {
2     println!("Hello, world!");
3 }
4
```

Design a set of component libraries that can be easily composed into large-scale FEM/BEM applications.

```
fn main() {  
H     let s1 = String::from("Hello"); => String      ■ move occurs because `s1` has type `String`  
H     let s2 = s1; => String      ■■■ value moved here  
E     println!("{}", s1);      ■■ borrow of moved value: `s1` value borrowed here after move  
}
```

- Every variable in Rust has a unique owner
- Assigning s1 to s2 transfers ownership of the string.

The Borrow Checker



```
fn main() {  
    let x = 5.0; => f64  
  
    let y = &x; => &f64  
    let z = &x; => &f64  
  
    println!("{} {}", y, z);  
}
```

- Creating references is called borrowing.
- We are not allowed to mutate a variable after it was borrowed.
- We can have either one mutable borrow or many non-mutable borrows, not both at the same time.

```
fn main() {  
    let mut x = 5.0; => f64  
  
H     let y = &x; => &f64      ■ borrow of `x` occurs here  
  
E     x = 2.0;      ■■ cannot assign to `x` because it is borrowed  assignment to borrowed `x` occurs here  
  
H     println!("This won't work {}", &y);      ■ borrow later used here  
}
```

What is already done?



Grids and finite elements

- Various order Lagrange elements on triangles
- Order 1 and order 2 basis function on quadrilaterals
- Degree 1 Raviart-Thomas functions on triangles
- Grids on single nodes and prototype multi-node grid management

FMM

- Distributed balanced octrees are implemented
- Single note experimental kernel independent FMM is working
- Currently working on distributed FMM loop

Integration

- Standard Gauss type rules for regular integrals over triangles and quadrilaterals
- Have fully numerical singular integrators on triangles (based on Erichsen/Sauter)
- Singular integrators on quadrilaterals almost done

Other work

- Developing core sparse linear algebra data structures
- Rust native dense matrix decompositions for randomized linear algebra
- SIMD Optimized kernel evaluation
- C API and Python interfaces

The Rust Scientific Computing ecosystem



Core utilities

- Num - Traits and utility routines for floating point types.
- Cauchy - A complex scalar type library (uses Num).
- approx - Unit test routines for floating point types.
- Serde - Serialization library.

Linear Algebra/Machine Learning

- ndarray - multidimensional array type in Rust.
- nalgebra - Rust linear algebra library with Lapack bindings.
- Linfa - Native machine learning in Rust.
- tensorflow - Rust bindings to tensorflow.
- petsc-rs - Rust bindings for Petsc.

Parallelization

- rsmpi - MPI Interface for Rust.
- Rayon - Fork/Join threading, parallel loops and iterators.
- async/await (Tokio, etc.) - frameworks for asynchronous execution.

Language bindings

- PyO3 - Binding complex Rust types to Python.
- maturin - Simple tools to wrap Rust C API bindings with Python cffi.
- Any language that supports C bindings can talk to Rust easily.

Still lots of moving parts in the ecosystem but an enormous amount of development happening.

- Using Python + JIT (e.g. Numba) works well for simple array oriented algorithms. More problematic for complex algorithms.
- We are at the start of a multi-year journey to develop very large-scale integral equation solvers in Rust.
- Rust is ready to be looked at by Scientific Computing Community (though far away from the level of maturity of the C++ ecosystem)
- Rust makes little sense when legacy C++ code exists or there is need for heterogeneous compute with SyCL or Cuda.
- But consider and evaluate Rust if you start a completely new project.